

Learning Feedback Fuzzers: AFL and libFuzzer



Dr. Jared DeMott

CTO AND FOUNDER

@jareddemott www.vdalabs.com



Overview



Feedback fuzzing

- Research
 - Genetic algorithm
 - Constraint solver
- Integrated unit testing
- libFuzzer
- AFL



Feedback Fuzzing Research

PathCrawler (2004)

DART (2004)

KLEE (2005-2006)

CUTE and jCUTE (2005-2006)


- Early tools were all Unit test tools for C
- jCUTE was Java

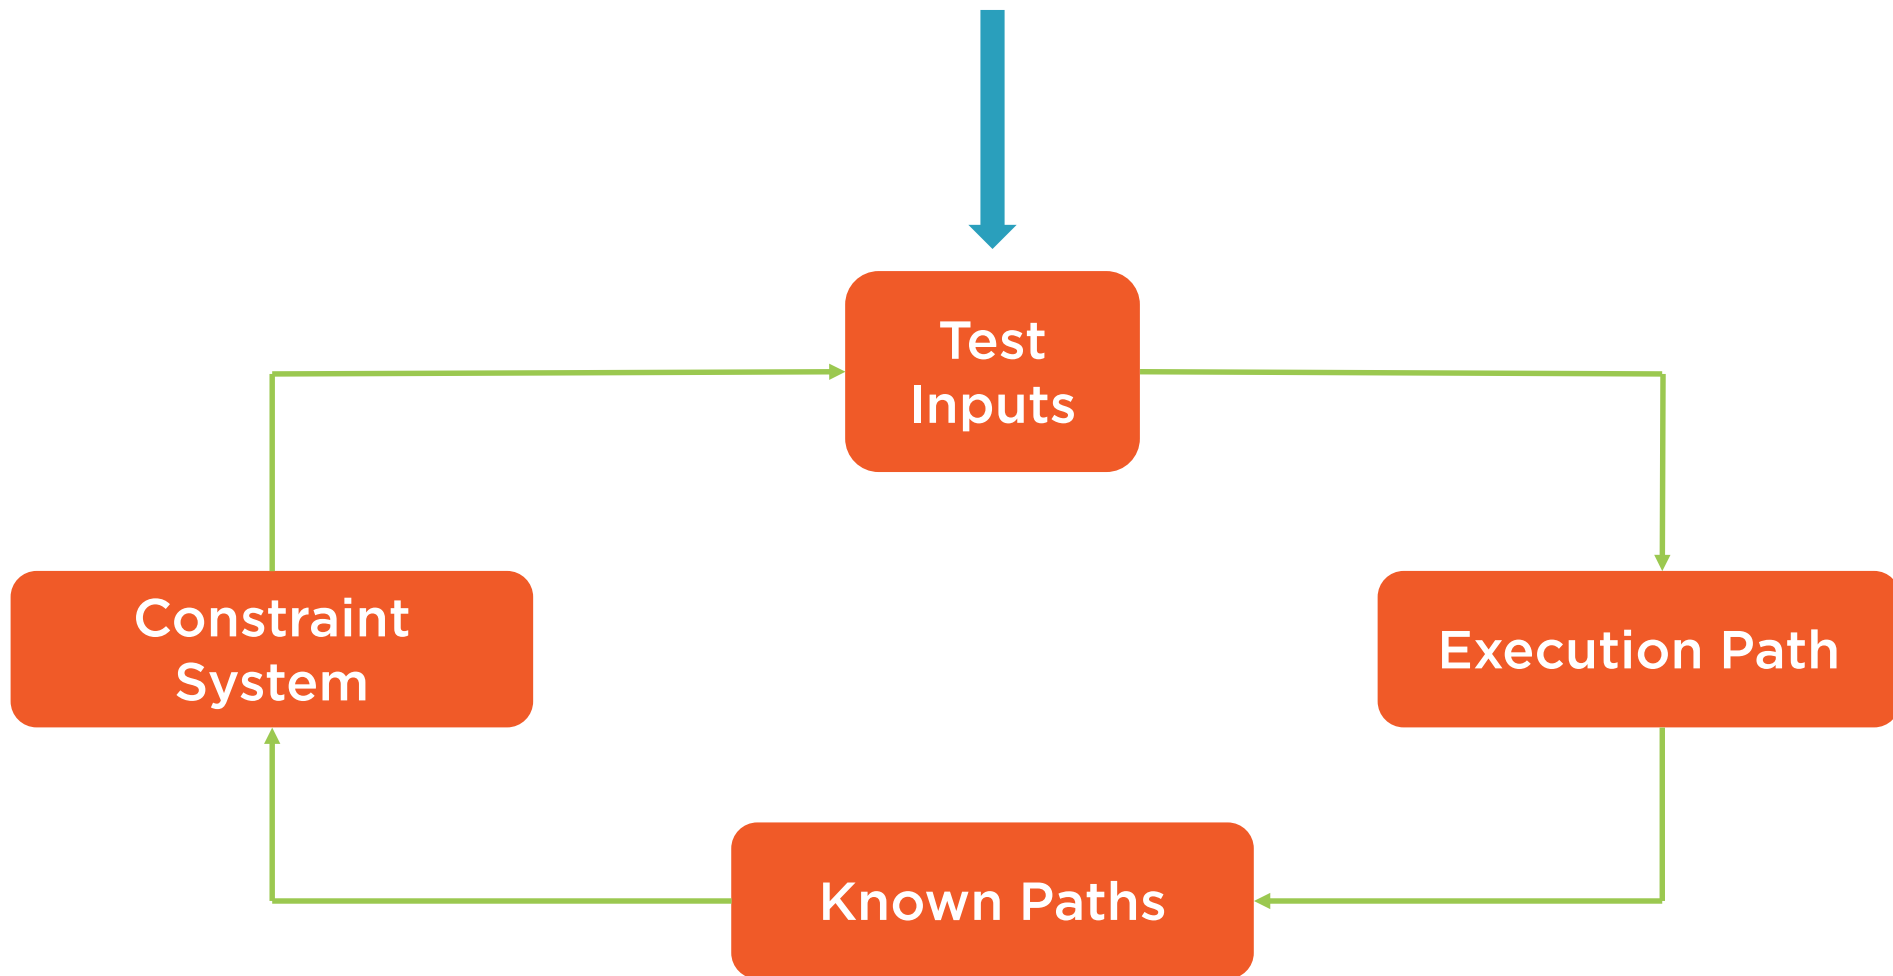
EFS (2007)

- First to use GA for fuzzing?

None of these were really fit for general use







Basic White-box Fuzzing Methodology



Concolic Testing

Symbolic execution + constraint solving +
dynamic testing == Concolic

- *A combination of whitebox and blackbox*
 - Reports on real issues like BB
- Finds more issues like WB
 - Increased accuracy
 - Can find the corner cases
 - Decreased runtime
 - Less need for continued test repeats or irrelevant tests



```
void f(int x, int y) {  
    int z = 2 * y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); // bug!  
        }  
    }  
}
```

Randomly reaching this code would be unlikely
A fuzzing sample that was close would be required
x = 100000, y = 50001



Limitations

Path Explosion

Program exhibits nondeterministic behavior

- May follow a different path than the intended one
- This can lead to non-termination of the search and poor coverage



Limitations

Even in a deterministic - issues

- Imprecise symbolic representations
 - Pointers, floating point numbers, etc.
- Incomplete theorem proving
- Failure to search most fruitful portion of a large or infinite path tree

Programs which thoroughly mix the state of their variables

- "if (md5_hash(input) == 0xdeadbeef)"



SAGE

Operates on native binaries

Found past “hard bugs”

- ANI-format bug that blackbox and whitebox missed
- MS07-017

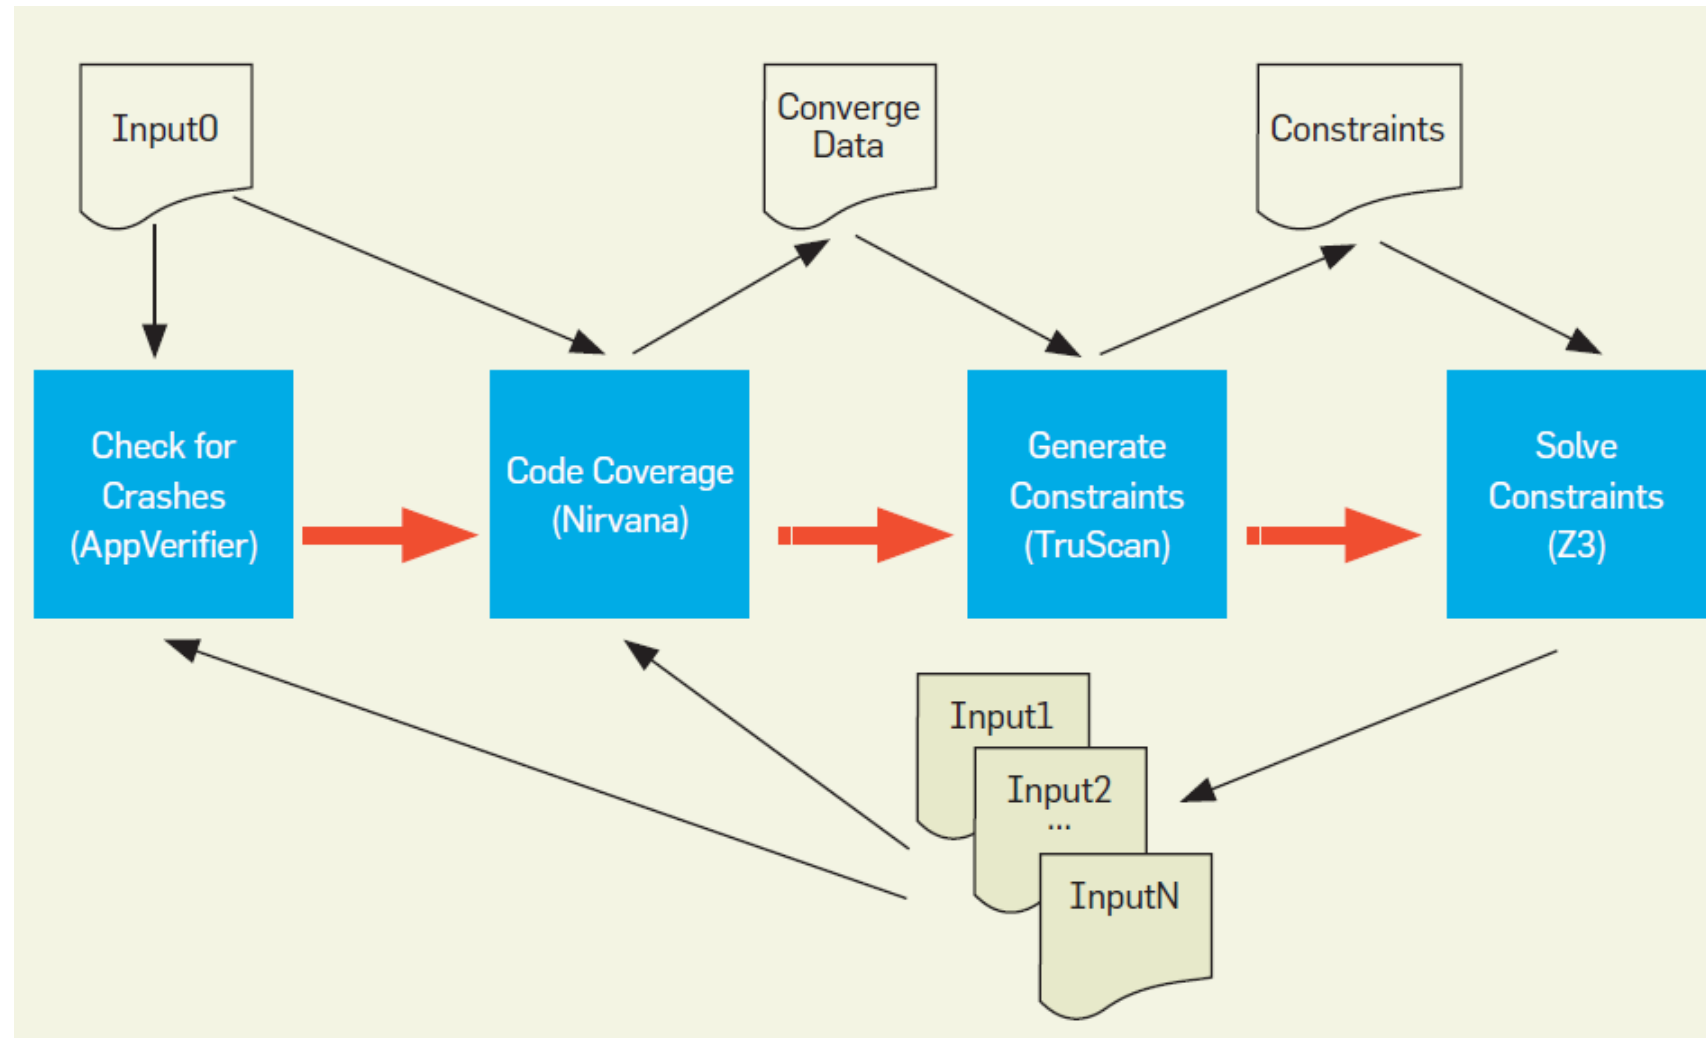
Found 1/3 of all file fuzzing bugs during the security testing of Windows 7

- SAGE is typically run last, meaning those bugs were missed by other approaches

Internal only tool

- Except for Springfield





Automatic Unit Tests

Pex

- Generates test suites with high code coverage using automated white box analysis
 - Works for .NET code
- Moles
 - Simulates interfaces, such as the file system or a database, so that unit testing can be conducted



Automatic Unit Tests

IntelliTest

- Explores.NET code to generate a suite of unit tests
 - For every statement, a test input is generated that will execute it
- <https://msdn.microsoft.com/en-us/library/dn823749.aspx>
- C# / 32 bit



libFuzzer

In-process, coverage-guided, evolutionary fuzzing engine

Linked with the library under test

- Feeds fuzzed inputs to the library via a specific fuzzing entrypoint
- Tracks which areas of the code are reached, and generates mutations on the corpus of input data in order to maximize the code coverage
- Code coverage information is provided by LLVM's SanitizerCoverage instrumentation



libFuzzer

Developer tool

- Build tools, write stub, compile, and run

Corpus

- GA search tools are much more efficient with a good set of seeds
- Samples that achieve new paths are added to this folder as they're discovered



libFuzzer

You need source

Run

- <http://llvm.org/docs/LibFuzzer.html>
- <https://github.com/google/fuzzer-test-suite/blob/master/tutorial/libFuzzerTutorial.md>



Demo



AFL fuzzer



American Fuzzy Lop

2014 released

- Compile-time instrumentation and genetic algorithms to discover clean, interesting test cases that trigger new internal states in the targeted binary
- <http://lcamtuf.coredump.cx/afl/>



```

7  bool FuzzMe(const uint8_t *Data, size_t DataSize) {
8      if( Data[0] == 'F')
9      {
10         if(Data[1] == 'U')
11         {
12             if(Data[2] == 'Z')
13             {
14                 if(Data[3] == 'Z')
15                 {
16                     abort(); // bug
17                 }
18             }
19         }
20     }
21
22     return 0;
23 }
24
25 int Test(const uint8_t *Data, size_t Size) {
26     FuzzMe(Data, Size);
27     return 0;
28 }
29
30 int main(int argc, char *argv[], char *envp[]) {
31     FILE *f;
32     char buf[4];
33     if( argc == 2) {
34
35         f = fopen(argv[1], "r");
36         if(f)
37         {
38             fread(buf, 1, 4, f);
39             Test((uint8_t *)buf, sizeof(buf));
40         }
41     }
42 }
43

```



```
[Jareds-MacBook-Pro-2:simple jareddemott$ ../../afl-clang fuzz_me.cc -o fuzz_me
afl-cc 2.35b by <lcamtuf@google.com>
afl-as 2.35b by <lcamtuf@google.com>
[+] Instrumented 23 locations (64-bit, non-hardened mode, ratio 100%).
Jareds-MacBook-Pro-2:simple jareddemott$ ../../afl-fuzz -i testcases/ -o findings/ -- ./fuzz_me @@
```



american fuzzy lop 2.35b (fuzz_me)

process timing

run time : 0 days, 0 hrs, 6 min, 32 sec
last new path : 0 days, 0 hrs, 1 min, 4 sec
last uniq crash : 0 days, 0 hrs, 1 min, 4 sec
last uniq hang : 0 days, 0 hrs, 3 min, 46 sec

cycle progress

now processing : 1 (25.00%)
paths timed out : 0 (0.00%)

stage progress

now trying : havoc
stage execs : 31/256 (12.11%)
total execs : 600k
exec speed : 1584/sec

fuzzing strategy yields

bit flips : 0/88, 0/84, 1/76
byte flips : 0/11, 0/7, 0/2
arithmetics : 0/612, 0/203, 0/68
known ints : 0/59, 0/138, 0/70
dictionary : 0/0, 0/0, 0/0
havoc : 2/571k, 0/27.5k
trim : 99.93%/19, 0.00%

overall results

cycles done : **684**
total paths : 4
uniq crashes : **1**
uniq hangs : 2

map coverage

map density : **0.01% / 0.02%**
count coverage : 1.00 bits/tuple

findings in depth

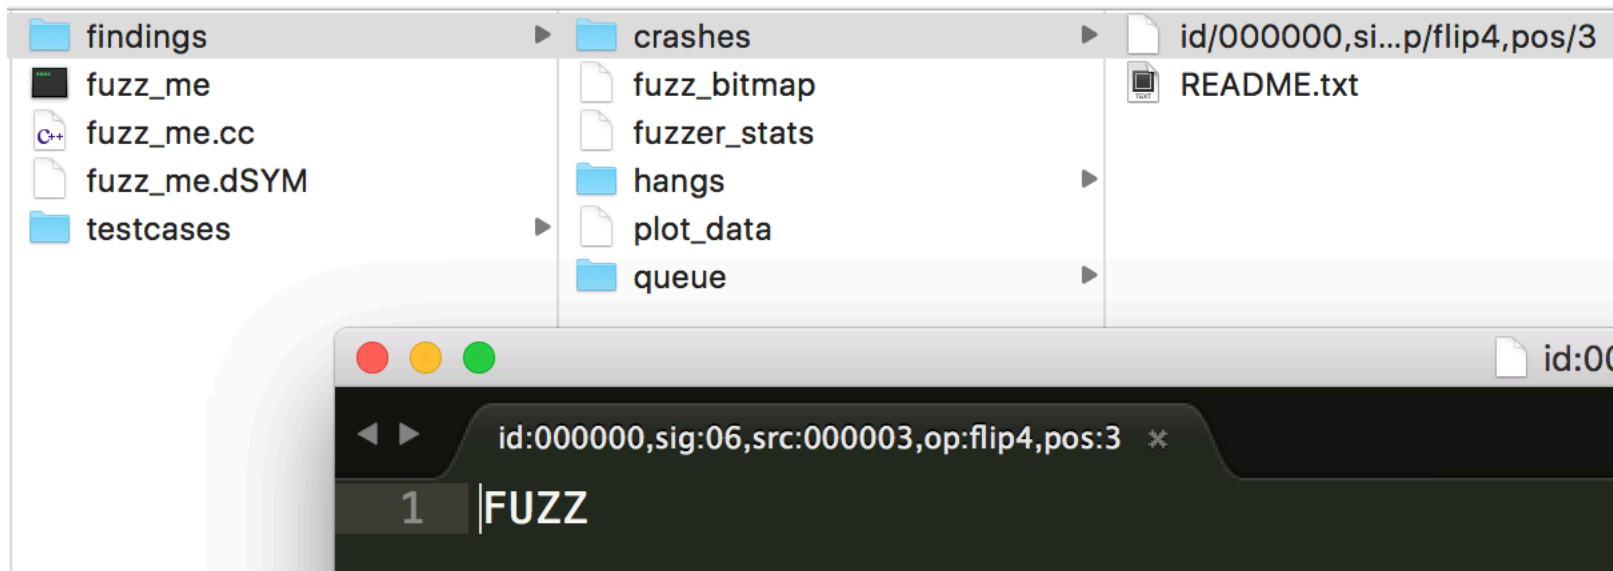
favorable paths : 4 (100.00%)
new edges on : 4 (100.00%)
total crashes : **15 (1 unique)**
total hangs : 6 (2 unique)

path geometry

levels : 3
pending : 0
pend fav : 0
own finds : 2
imported : n/a
stability : 100.00%

[cpu: **42%**]





Why It Works?

Not just the feedback

- Practical engineering
 - Uses fuzzing best practices to find new paths quickly
 - Not many knobs
 - Includes helpers like minset
 - Monitoring, coverage, distribution – all in one tool



AFL

Limitations:

- You mostly need source
 - There is a Linux/QEMU combo that allows black-box fuzzing
- Not every program is suitable



Summary



Closed

.NET

Open

- libFuzzer
- AFL
- GA works well in practice, but research still continues around constraint solving

Finish with metrics